

Practical Machine Learning Prediction Project

Removed For Privacy [GitHub](#)

Contents

Prepare the datasets	1
Train a prediction model	3
Generate the files for submission.	6

```
## Run time: 2015-08-23 18:38:29
## R version: R version 3.1.2 (2014-10-31)
```

Prepare the datasets

```
library(rpart)
library(rpart.plot)
```

```
## Warning: package 'rpart.plot' was built under R version 3.1.3
```

```
library(RColorBrewer)
```

```
## Warning: package 'RColorBrewer' was built under R version 3.1.3
```

```
library(rattle)
```

```
## Warning: package 'rattle' was built under R version 3.1.3
```

```
## Loading required package: RGtk2
```

```
## Warning: package 'RGtk2' was built under R version 3.1.3
```

```
## Rattle: A free graphical interface for data mining with R.
## Version 3.5.0 Copyright (c) 2006-2015 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.1.3
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 3.1.3
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.1.3
```

```
## randomForest 4.6-10
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

Load the training and testing data into a data table.

```
trainUrl <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
testUrl <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"
training <- read.csv(url(trainUrl), na.strings=c("NA", "#DIV/0!", ""))
testing <- read.csv(url(testUrl), na.strings=c("NA", "#DIV/0!", ""))
```

We are given an overwhelmingly large training data. We simply partition and generate a testing set rather than use only 20 observation in the given test set. Partition into training and testing is done with a 60/40 split in the code below.

```
inTrain <- createDataPartition(y=training$classe, p=0.6, list=FALSE)
myTraining <- training[inTrain, ]; myTesting <- training[-inTrain, ]
dim(myTraining); dim(myTesting)
```

```
## [1] 11776 160
```

```
## [1] 7846 160
```

Remove zero variance predictors

```
myDataNZV <- nearZeroVar(myTraining, saveMetrics=TRUE)

myNZVvars <- names(myTraining) %in% row.names(myDataNZV[myDataNZV$nzv==TRUE,])
myTraining <- myTraining[!myNZVvars]
#To check the new N?? of observations
dim(myTraining)
```

```
## [1] 11776 128
```

We remove the first column of the data and also clean Variables with too many NAs. For Variables that have more than a 60% threshold of NA's we leave them out.

```
myTraining <- myTraining[c(-1)]
trainingV3 <- myTraining #creating another subset to iterate in loop
for(i in 1:length(myTraining)) { #for every column in the training dataset
  if( sum( is.na( myTraining[, i] ) ) /nrow(myTraining) >= .6 ) { #if n?? NAs > 60% of total obser
    for(j in 1:length(trainingV3)) {
      if( length( grep(names(myTraining[i]), names(trainingV3)[j]) ) ==1) { #if the columns are
        trainingV3 <- trainingV3[ , -j] #Remove that column
      }
    }
  }
}
#To check the new N?? of observations
dim(trainingV3)
```

```
## [1] 11776    58
```

```
#Setting back to our set:  
myTraining <- trainingV3  
rm(trainingV3)
```

Do the same for the testing set of variables (as well as the original testing set.)

```
clean1 <- colnames(myTraining)  
clean2 <- colnames(myTraining[, -58]) #already with classe column removed  
myTesting <- myTesting[clean1]  
testing <- testing[clean2]  
  
#To check the new N?? of observations  
dim(myTesting)
```

```
## [1] 7846    58
```

```
dim(testing)
```

```
## [1] 20 57
```

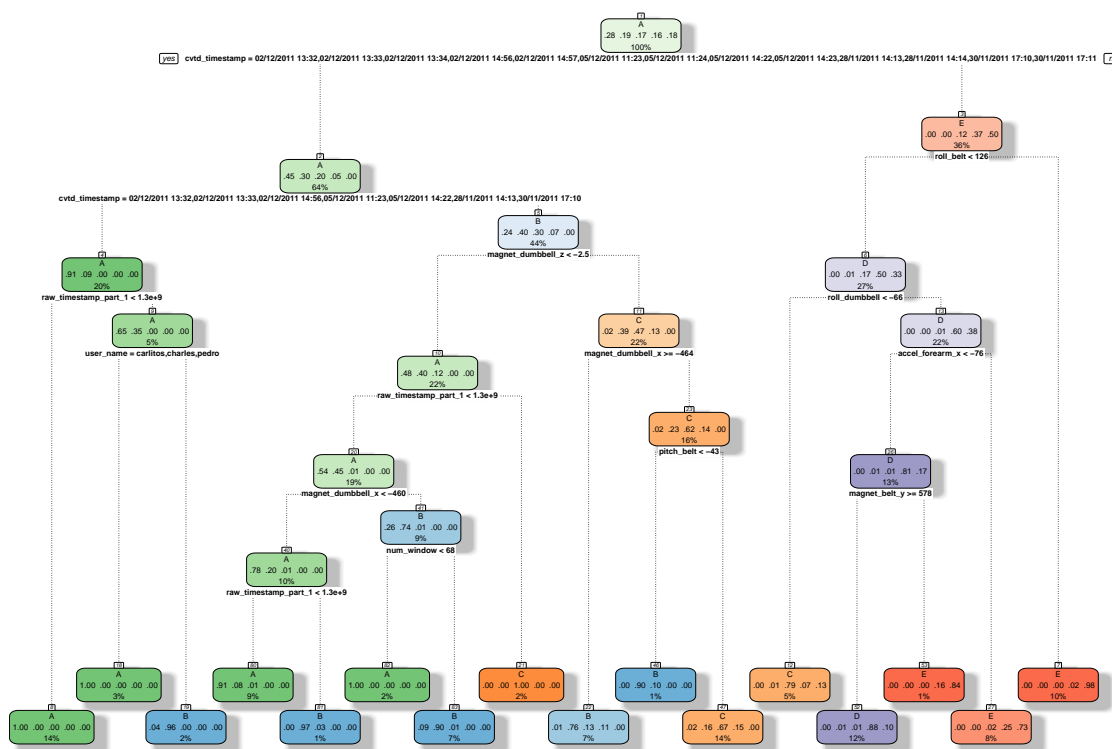
Harmonize the type of test and training data. Otherwise RandomForest throws an error that “Predictors in the new data do not match that of the training data”

```
for (i in 1:length(testing) ) {  
  for(j in 1:length(myTraining)) {  
    if( length( grep(names(myTraining[i]), names(testing)[j]) ) ==1) {  
      class(testing[j]) <- class(myTraining[i])  
    }  
  }  
}  
  
#And to make sure Coertion really worked, simple smart ass technique:  
testing <- rbind(myTraining[2, -58] , testing) #note row 2 does not mean anything, this will be removed  
testing <- testing[-1,]
```

Train a prediction model

We first use decision trees.

```
modFitA1 <- rpart(classe ~ ., data=myTraining, method="class")  
fancyRpartPlot(modFitA1)
```



Rattle 2015–Aug-23 18:39:04 siva

We check how good the prediction is

```
predictionsA1 <- predict(modFitA1, myTesting, type = "class")
```

We look at the cross-tabulation of observed and predicted classes with associated statistics.

```
confusionMatrix(predictionsA1, myTesting$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 2143   64    7    3    0
##           B   69 1269   77   70    0
##           C   20  174 1263  198   63
##           D    0   11   12  823   80
##           E    0    0    9  192 1299
##
## Overall Statistics
##
##           Accuracy : 0.8663
##           95% CI : (0.8586, 0.8738)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
```

```
## Kappa : 0.8308
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
## Class: A Class: B Class: C Class: D Class: E
## Sensitivity 0.9601 0.8360 0.9232 0.6400 0.9008
## Specificity 0.9868 0.9659 0.9298 0.9843 0.9686
## Pos Pred Value 0.9666 0.8545 0.7352 0.8888 0.8660
## Neg Pred Value 0.9842 0.9609 0.9829 0.9331 0.9775
## Prevalence 0.2845 0.1935 0.1744 0.1639 0.1838
## Detection Rate 0.2731 0.1617 0.1610 0.1049 0.1656
## Detection Prevalence 0.2826 0.1893 0.2190 0.1180 0.1912
## Balanced Accuracy 0.9735 0.9009 0.9265 0.8121 0.9347
```

Since accuracy is poor we try to use RandomForest

```
modFitB1 <- randomForest(classe ~. , data=myTraining)
predictionsB1 <- predict(modFitB1, myTesting, type = "class")
confusionMatrix(predictionsB1, myTesting$classe)
```

```
## Confusion Matrix and Statistics
##
## Reference
## Prediction A B C D E
## A 2232 1 0 0 0
## B 0 1517 0 0 0
## C 0 0 1367 5 0
## D 0 0 1 1281 2
## E 0 0 0 0 1440
##
## Overall Statistics
##
## Accuracy : 0.9989
## 95% CI : (0.9978, 0.9995)
## No Information Rate : 0.2845
## P-Value [Acc > NIR] : < 2.2e-16
##
## Kappa : 0.9985
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
## Class: A Class: B Class: C Class: D Class: E
## Sensitivity 1.0000 0.9993 0.9993 0.9961 0.9986
## Specificity 0.9998 1.0000 0.9992 0.9995 1.0000
## Pos Pred Value 0.9996 1.0000 0.9964 0.9977 1.0000
## Neg Pred Value 1.0000 0.9998 0.9998 0.9992 0.9997
## Prevalence 0.2845 0.1935 0.1744 0.1639 0.1838
## Detection Rate 0.2845 0.1933 0.1742 0.1633 0.1835
## Detection Prevalence 0.2846 0.1933 0.1749 0.1637 0.1835
## Balanced Accuracy 0.9999 0.9997 0.9992 0.9978 0.9993
```

We obtain much better accuracy and pick this model.

Generate the files for submission.

We use the random forest model to generate the required files.

```
predictionsB2 <- predict(modFitB1, testing, type = "class")

pml_write_files = function(x){
  n = length(x)
  for(i in 1:n){
    filename = paste0("problem_id_",i,".txt")
    write.table(x[i],file=filename,quote=FALSE,row.names=FALSE,col.names=FALSE)
  }
}

pml_write_files(predictionsB2)
```